

# Gateway Service Deployment Guide

Production deployment guide for the Mission Engadi Gateway Service

## Table of Contents

1. [Overview](#)
2. [Prerequisites](#)
3. [Local Development](#)
4. [Docker Deployment](#)
5. [Kubernetes Deployment](#)
6. [Cloud Platform Deployment](#)
7. [Configuration](#)
8. [Monitoring & Observability](#)
9. [Scaling](#)
10. [Troubleshooting](#)

## Overview

This guide covers deploying the Gateway Service in various environments, from local development to production cloud deployments.

## Deployment Options

- **Local Development:** Docker Compose
- **Container Orchestration:** Kubernetes, Docker Swarm
- **Cloud Platforms:** AWS, GCP, Azure, Fly.io
- **Serverless:** AWS Lambda, Google Cloud Functions (with adapters)

## Prerequisites

### Required

- **Python:** 3.11 or higher
- **PostgreSQL:** 15 or higher
- **Docker:** 20.10 or higher (for containerized deployment)
- **Git:** For version control

### Recommended

- **Redis:** 7 or higher (for caching and rate limiting)
- **Load Balancer:** For production deployments
- **Monitoring:** Prometheus, Grafana, or cloud-native solutions
- **Log Aggregation:** ELK stack, CloudWatch, or similar

# Local Development

## Quick Start with Scripts

```
# Clone repository
git clone https://github.com/mission-engadi/gateway_service.git
cd gateway_service

# Start service
./start.sh

# Check status
./status.sh

# View logs
tail -f unicorn.log

# Stop service
./stop.sh

# Restart service
./restart.sh
```

## Manual Setup

### 1. Create virtual environment

```
bash
python3 -m venv venv
source venv/bin/activate
```

### 2. Install dependencies

```
bash
pip install -r requirements.txt
```

### 3. Configure environment

```
bash
cp .env.example .env
# Edit .env with your configuration
```

### 4. Start PostgreSQL

```
bash
docker run -d \
--name gateway-postgres \
-p 5432:5432 \
-e POSTGRES_PASSWORD=postgres \
-e POSTGRES_DB=gateway_db \
postgres:15-alpine
```

### 5. Run migrations

```
bash
alembic upgrade head
```

## 6. Start service

```
bash
uvicorn app.main:app --reload --port 8000
```

---

# Docker Deployment

## Using Docker Compose

### docker-compose.yml

```
version: '3.8'

services:
  gateway:
    build: .
    ports:
      - "8000:8000"
    environment:
      - DATABASE_URL=postgresql+asyncpg://postgres:postgres@postgres:5432/gateway_db
      - REDIS_URL=redis://redis:6379/0
      - SECRET_KEY=${SECRET_KEY}
      - ENVIRONMENT=production
    depends_on:
      - postgres
      - redis
    restart: unless-stopped
    healthcheck:
      test: ["CMD", "curl", "-f", "http://localhost:8000/api/v1/health"]
      interval: 30s
      timeout: 10s
      retries: 3

  postgres:
    image: postgres:15-alpine
    environment:
      - POSTGRES_PASSWORD=postgres
      - POSTGRES_DB=gateway_db
    volumes:
      - postgres_data:/var/lib/postgresql/data
    restart: unless-stopped

  redis:
    image: redis:7-alpine
    restart: unless-stopped

volumes:
  postgres_data:
```

## Deployment Commands

```
# Build and start
docker-compose up -d --build

# View logs
docker-compose logs -f gateway

# Scale gateway instances
docker-compose up -d --scale gateway=3

# Stop services
docker-compose down

# Remove volumes (CAREFUL!)
docker-compose down -v
```

## Using Docker Directly

```
# Build image
docker build -t gateway-service:latest .

# Run container
docker run -d \
  --name gateway \
  -p 8000:8000 \
  -e DATABASE_URL=postgresql://... \
  -e SECRET_KEY=your-secret-key \
  gateway-service:latest

# View logs
docker logs -f gateway

# Stop container
docker stop gateway

# Remove container
docker rm gateway
```

# Kubernetes Deployment

## Deployment Manifests

### deployment.yaml

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: gateway-service
  labels:
    app: gateway-service
spec:
  replicas: 3
  selector:
    matchLabels:
      app: gateway-service
  template:
    metadata:
      labels:
        app: gateway-service
    spec:
      containers:
        - name: gateway
          image: mission-engadi/gateway-service:latest
          ports:
            - containerPort: 8000
          env:
            - name: DATABASE_URL
              valueFrom:
                secretKeyRef:
                  name: gateway-secrets
                  key: database-url
            - name: SECRET_KEY
              valueFrom:
                secretKeyRef:
                  name: gateway-secrets
                  key: secret-key
            - name: ENVIRONMENT
              value: "production"
      resources:
        requests:
          memory: "256Mi"
          cpu: "250m"
        limits:
          memory: "512Mi"
          cpu: "500m"
      livenessProbe:
        httpGet:
          path: /api/v1/live
          port: 8000
        initialDelaySeconds: 30
        periodSeconds: 10
      readinessProbe:
        httpGet:
          path: /api/v1/ready
          port: 8000
        initialDelaySeconds: 5
        periodSeconds: 5

```

## service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: gateway-service
spec:
  selector:
    app: gateway-service
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8000
  type: LoadBalancer
```

## secrets.yaml

```
apiVersion: v1
kind: Secret
metadata:
  name: gateway-secrets
type: Opaque
stringData:
  database-url: "postgresql+asyncpg://user:password@postgres:5432/gateway_db"
  secret-key: "your-secret-key-here"
```

## ingress.yaml

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: gateway-ingress
  annotations:
    cert-manager.io/cluster-issuer: "letsencrypt-prod"
spec:
  tls:
    - hosts:
        - api.engadi.org
      secretName: gateway-tls
  rules:
    - host: api.engadi.org
      http:
        paths:
          - path: /
            pathType: Prefix
            backend:
              service:
                name: gateway-service
                port:
                  number: 80
```

## Deployment Commands

```
# Create namespace
kubectl create namespace mission-engadi

# Apply secrets
kubectl apply -f secrets.yaml -n mission-engadi

# Deploy application
kubectl apply -f deployment.yaml -n mission-engadi
kubectl apply -f service.yaml -n mission-engadi
kubectl apply -f ingress.yaml -n mission-engadi

# Check deployment
kubectl get pods -n mission-engadi
kubectl get svc -n mission-engadi

# View logs
kubectl logs -f deployment/gateway-service -n mission-engadi

# Scale deployment
kubectl scale deployment gateway-service --replicas=5 -n mission-engadi

# Update deployment
kubectl set image deployment/gateway-service \
  gateway=mission-engadi/gateway-service:v1.1.0 \
  -n mission-engadi

# Rollback deployment
kubectl rollout undo deployment/gateway-service -n mission-engadi
```

# Cloud Platform Deployment

---

## AWS (Elastic Beanstalk)

```
# Install EB CLI
pip install awsebcli

# Initialize EB application
eb init -p python-3.11 gateway-service

# Create environment
eb create gateway-prod

# Set environment variables
eb setenv \
  SECRET_KEY=your-secret-key \
  DATABASE_URL=postgresql://... \
  ENVIRONMENT=production

# Deploy
eb deploy

# Open in browser
eb open

# View logs
eb logs
```

## Google Cloud Platform (Cloud Run)

```
# Build and push to Container Registry
gcloud builds submit --tag gcr.io/PROJECT_ID/gateway-service

# Deploy to Cloud Run
gcloud run deploy gateway-service \
  --image gcr.io/PROJECT_ID/gateway-service \
  --platform managed \
  --region us-central1 \
  --allow-unauthenticated \
  --set-env-vars="SECRET_KEY=your-key,DATABASE_URL=postgresql://..."

# View logs
gcloud run logs read gateway-service
```

## Azure (Container Instances)

```
# Login to Azure
az login

# Create resource group
az group create --name gateway-rg --location eastus

# Create container instance
az container create \
--resource-group gateway-rg \
--name gateway-service \
--image mission-engadi/gateway-service:latest \
--dns-name-label gateway-engadi \
--ports 8000 \
--environment-variables \
SECRET_KEY=your-key \
DATABASE_URL=postgresql://...

# View logs
az container logs --resource-group gateway-rg --name gateway-service
```

## Fly.io

```
# Install flyctl
curl -L https://fly.io/install.sh | sh

# Login
fly auth login

# Launch app
fly launch --name gateway-service

# Set secrets
fly secrets set \
SECRET_KEY=your-secret-key \
DATABASE_URL=postgresql://...

# Deploy
fly deploy

# Scale
fly scale count 3

# View logs
fly logs
```

# Configuration

## Environment Variables

### Required

```
# Application
SECRET_KEY="your-secret-key-here" # Generate with: openssl rand -hex 32
DATABASE_URL="postgresql+asyncpg://user:pass@host:5432/db"

# Environment
ENVIRONMENT="production" # production, staging, development
DEBUG="false"
```

### Optional

```
# Redis
REDIS_URL="redis://localhost:6379/0"

# CORS
CORS_ORIGINS="https://app.engadi.org,https://engadi.org"

# Rate Limiting
DEFAULT_RATE_LIMIT="100"
DEFAULT_RATE_WINDOW="60"

# Monitoring
DATADOG_API_KEY="your-datadog-key"
SENTRY_DSN="your-sentry-dsn"

# Logging
LOG_LEVEL="INFO" # DEBUG, INFO, WARNING, ERROR, CRITICAL
```

## Database Setup

### PostgreSQL

```
-- Create database
CREATE DATABASE gateway_db;

-- Create user
CREATE USER gateway_user WITH PASSWORD 'secure_password';

-- Grant permissions
GRANT ALL PRIVILEGES ON DATABASE gateway_db TO gateway_user;
```

## Run Migrations

```
# Upgrade to latest version
alembic upgrade head

# Check current version
alembic current

# Downgrade if needed
alembic downgrade -1
```

# Monitoring & Observability

## Health Checks

```
# Liveness probe
curl http://localhost:8000/api/v1/live

# Readiness probe
curl http://localhost:8000/api/v1/ready

# Health check
curl http://localhost:8000/api/v1/health
```

## Logging

### Structured Logging

Production logs are in JSON format:

```
{
  "timestamp": "2024-12-25T00:00:00Z",
  "level": "INFO",
  "logger": "app.services.proxy",
  "message": "Request proxied successfully",
  "trace_id": "abc123",
  "user_id": "user456",
  "status_code": 200,
  "response_time": 0.15
}
```

### Log Aggregation

#### Docker Compose:

```
docker-compose logs -f gateway
```

#### Kubernetes:

```
kubectl logs -f deployment/gateway-service -n mission-engadi
```

#### Cloud Platforms:

- AWS: CloudWatch Logs
- GCP: Cloud Logging
- Azure: Azure Monitor

## Metrics

### Prometheus Integration

```
# prometheus.yml
scrape_configs:
  - job_name: 'gateway'
    static_configs:
      - targets: ['gateway:8000']
```

### Key Metrics

- **Request rate:** requests per second
- **Response time:** p50, p95, p99 latencies
- **Error rate:** 4xx and 5xx responses
- **Active connections:** current connections
- **Service health:** backend service status

## Scaling

### Horizontal Scaling

#### Docker Compose

```
docker-compose up -d --scale gateway=5
```

#### Kubernetes

```
# Manual scaling
kubectl scale deployment gateway-service --replicas=5

# Auto-scaling
kubectl autoscale deployment gateway-service \
  --cpu-percent=70 \
  --min=3 \
  --max=10
```

### Vertical Scaling

Update resource limits:

```
resources:
  requests:
    memory: "512Mi"
    cpu: "500m"
  limits:
    memory: "1Gi"
    cpu: "1000m"
```

## Load Balancing

- **Docker Compose:** Use Nginx or HAProxy
- **Kubernetes:** Built-in service load balancing

- **Cloud:** Use cloud load balancers (ALB, Cloud Load Balancing, etc.)
- 

## Troubleshooting

### Service Won't Start

**1. Check logs:**

```
bash
docker logs gateway
kubectl logs deployment/gateway-service
```

**2. Verify environment variables:**

```
bash
docker exec gateway env
```

**3. Test database connection:**

```
bash
docker exec gateway python -c "from app.db.session import engine; print('Connected')"
```

### High Response Times

**1. Check resource usage:**

```
bash
docker stats gateway
kubectl top pods
```

**2. Review slow queries:**

```
bash
# Check PostgreSQL slow query log
docker exec postgres psql -U postgres -c "SELECT * FROM pg_stat_activity WHERE state = 'active';"
```

**3. Scale horizontally:**

```
bash
docker-compose up -d --scale gateway=3
```

### Database Connection Issues

**1. Verify DATABASE\_URL:**

```
bash
echo $DATABASE_URL
```

**2. Test connection:**

```
bash
psql $DATABASE_URL
```

**3. Check connection pool:**

- Increase pool size in config
- Check for connection leaks

## Memory Leaks

### 1. Monitor memory usage:

```
bash
  docker stats gateway
```

### 2. Enable memory profiling:

```
python
# In app/main.py
import tracemalloc
tracemalloc.start()
```

### 3. Restart service periodically (temporary fix)

---

## Production Checklist

### Security

- [ ] Strong SECRET\_KEY generated
- [ ] HTTPS enabled (TLS certificates)
- [ ] Database credentials secured
- [ ] Environment variables not in code
- [ ] CORS origins configured
- [ ] Rate limiting enabled
- [ ] Security headers configured

### Performance

- [ ] Database connection pooling configured
- [ ] Redis caching enabled
- [ ] CDN configured for static assets
- [ ] Compression enabled
- [ ] Health checks configured

### Monitoring

- [ ] Logging configured
- [ ] Metrics collection enabled
- [ ] Alerts configured
- [ ] Health check endpoints tested
- [ ] Error tracking enabled (Sentry)

### Backup & Recovery

- [ ] Database backups automated
  - [ ] Disaster recovery plan documented
  - [ ] Rollback procedure tested
  - [ ] Data retention policy defined
-

## Support

---

- **Documentation:** [docs.engadi.org](https://docs.engadi.org) (<https://docs.engadi.org>)
- **Issues:** [GitHub Issues](https://github.com/mission-engadi/gateway_service/issues) ([https://github.com/mission-engadi/gateway\\_service/issues](https://github.com/mission-engadi/gateway_service/issues))
- **Email:** support@engadi.org