

# Gateway Service Integration Guide

How to integrate your microservice with the Mission Engadi Gateway

## Table of Contents

1. [Overview](#)
2. [Quick Start](#)
3. [Registering Your Service](#)
4. [Authentication](#)
5. [Rate Limiting](#)
6. [Health Checks](#)
7. [Error Handling](#)
8. [Request Tracing](#)
9. [Best Practices](#)
10. [Examples](#)

## Overview

The Gateway Service acts as the single entry point for all Mission Engadi microservices. It provides:

- **Unified routing:** All requests go through the gateway
- **Authentication:** Centralized JWT validation
- **Rate limiting:** Protect your service from overload
- **Health monitoring:** Automatic health checks
- **Request tracing:** End-to-end request tracking
- **Logging:** Centralized request/response logging

## Architecture Diagram

```
Client → Gateway Service → Your Microservice
      |   +-- Routing
      |   +-- Auth
      |   +-- Rate Limiting
      |   +-- Health Checks
      |   +-- Logging
```

## Quick Start

### 1. Deploy Your Service

Ensure your service is accessible from the gateway:

```
# docker-compose.yml
services:
  your-service:
    build: .
    ports:
      - "8003:8003"
    networks:
      - mission-engadi

networks:
  mission-engadi:
    external: true
```

## 2. Register with Gateway

Create a route configuration:

```
curl -X POST http://gateway:8000/api/v1/routes \
-H "Content-Type: application/json" \
-H "Authorization: Bearer <admin-token>" \
-d '{
  "path": "/api/v1/your-service/*",
  "target_service": "your-service",
  "target_url": "http://your-service:8003",
  "methods": ["GET", "POST", "PUT", "DELETE"],
  "auth_required": true,
  "rate_limit": 100,
  "is_active": true
}'
```

## 3. Test Your Integration

```
# Request goes through gateway
curl http://gateway:8000/api/v1/your-service/health

# Gateway forwards to: http://your-service:8003/api/v1/your-service/health
```

# Registering Your Service

## Route Configuration

### Path Patterns

Supports wildcard matching:

```
# Exact match
"/api/v1/auth/login" # Only matches this exact path

# Wildcard match
"/api/v1/content/**" # Matches /api/v1/content/items, /api/v1/content/items/123, etc.

# Multiple wildcards
"/api/v1/**/items/**" # Matches /api/v1/content/items/123, /api/v1/media/items/456
```

## HTTP Methods

Specify allowed methods:

```
{
  "methods": ["GET", "POST"], // Only GET and POST allowed
  // or
  "methods": ["*"], // All methods allowed
}
```

## Authentication

Control authentication requirements:

```
{
  "auth_required": true, // Requires JWT token
  "auth_required": false, // Public endpoint
}
```

## Using the API

### Create Route (Admin Only)

```
import requests

response = requests.post(
    "http://gateway:8000/api/v1/routes",
    headers={
        "Authorization": f"Bearer {admin_token}",
        "Content-Type": "application/json"
    },
    json={
        "path": "/api/v1/media/*",
        "target_service": "media-service",
        "target_url": "http://media-service:8004",
        "methods": ["GET", "POST", "PUT", "DELETE"],
        "auth_required": True,
        "rate_limit": 50,
        "is_active": True
    }
)

route = response.json()
print(f"Route created with ID: {route['id']}")
```

### Update Route

```
route_id = 123

response = requests.put(
    f"http://gateway:8000/api/v1/routes/{route_id}",
    headers={"Authorization": f"Bearer {admin_token}"},
    json={
        "rate_limit": 100, # Increase rate limit
        "is_active": True
    }
)
```

## Delete Route

```
response = requests.delete(
    f"http://gateway:8000/api/v1/routes/{route_id}",
    headers={"Authorization": f"Bearer {admin_token}"}
)
```

## Authentication

### How It Works

1. Client sends request with JWT token to gateway
2. Gateway validates token (if `auth_required: true`)
3. Gateway forwards request to your service
4. Your service receives validated request

### Accessing User Information

The gateway adds authentication headers:

```
# In your service
from fastapi import Request

@app.get("/api/v1/your-service/profile")
async def get_profile(request: Request):
    # Gateway adds these headers after validation
    user_id = request.headers.get("X-Gateway-User-ID")
    user_email = request.headers.get("X-Gateway-User-Email")
    user_roles = request.headers.get("X-Gateway-User-Roles").split(", ")

    return {
        "user_id": user_id,
        "email": user_email,
        "roles": user_roles
    }
```

### Headers Added by Gateway

```
X-Gateway-User-ID: user_123
X-Gateway-User-Email: user@example.com
X-Gateway-User-Roles: user,editor
X-Gateway-Request-ID: abc123-def456
X-Forwarded-For: 192.168.1.100
```

### Public Endpoints

For public endpoints, set `auth_required: false`:

```
{
    "path": "/api/v1/public/*",
    "auth_required": false
}
```

# Rate Limiting

## Default Rate Limits

Set at route level:

```
{
  "path": "/api/v1/your-service/*",
  "rate_limit": 100 // 100 requests per minute
}
```

## Custom Rate Limit Rules

Create specific rules for endpoints:

```
# Strict rate limit for login endpoint
requests.post(
    "http://gateway:8000/api/v1/rate-limits",
    headers={"Authorization": f"Bearer {admin_token}"},
    json={
        "path": "/api/v1/your-service/login",
        "limit": 5,      # 5 requests
        "window": 60,    # per 60 seconds
        "is_active": True
    }
)

# Generous limit for read endpoints
requests.post(
    "http://gateway:8000/api/v1/rate-limits",
    headers={"Authorization": f"Bearer {admin_token}"},
    json={
        "path": "/api/v1/your-service/items",
        "limit": 1000,
        "window": 60,
        "is_active": True
    }
)
```

## Rate Limit Response

When rate limit is exceeded:

```
{
  "error": {
    "code": "RATE_LIMIT_EXCEEDED",
    "message": "Too many requests",
    "retry_after": 45
  }
}
```

Headers:

```
X-RateLimit-Limit: 100
X-RateLimit-Remaining: 0
X-RateLimit-Reset: 1640390400
Retry-After: 45
```

## Health Checks

### Implementing Health Endpoint

Your service must implement a health check endpoint:

```
from fastapi import FastAPI

app = FastAPI()

@app.get("/api/v1/health")
async def health_check():
    # Check database, Redis, etc.
    return {
        "status": "healthy",
        "service": "your-service",
        "version": "1.0.0",
        "checks": {
            "database": "connected",
            "redis": "connected"
        }
    }
```

### Health Check Intervals

Gateway checks service health:

- **Interval:** Every 30 seconds
- **Timeout:** 5 seconds
- **Failure threshold:** 3 consecutive failures

### Circuit Breaker

When your service is unhealthy:

1. Gateway marks service as “unhealthy”
2. Requests receive `503 Service Unavailable`
3. Gateway continues health checks
4. Service auto-recovers when healthy

## Error Handling

### Error Response Format

Your service should return consistent error responses:

```

from fastapi import HTTPException

@app.get("/api/v1/your-service/items/{item_id}")
async def get_item(item_id: int):
    item = await get_item_from_db(item_id)

    if not item:
        raise HTTPException(
            status_code=404,
            detail={
                "code": "ITEM_NOT_FOUND",
                "message": f"Item {item_id} not found",
                "details": {"item_id": item_id}
            }
        )

    return item

```

## Gateway Error Responses

### 502 Bad Gateway

Your service returned an invalid response

### 503 Service Unavailable

Your service is marked unhealthy

### 504 Gateway Timeout

Your service didn't respond within timeout (30 seconds)

## Request Tracing

### Trace ID Propagation

Gateway adds trace ID to all requests:

```

from fastapi import Request
import logging

@app.get("/api/v1/your-service/items")
async def get_items(request: Request):
    # Get trace ID from gateway
    trace_id = request.headers.get("X-Gateway-Request-ID")

    # Add to logs
    logging.info(
        "Fetching items",
        extra={"trace_id": trace_id}
    )

    # Include in responses for debugging
    return {
        "items": [...],
        "trace_id": trace_id
    }

```

## Structured Logging

Use structured logging for better observability:

```
import structlog

logger = structlog.get_logger()

@app.post("/api/v1/your-service/items")
async def create_item(item: ItemCreate, request: Request):
    trace_id = request.headers.get("X-Gateway-Request-ID")

    logger.info(
        "item.create",
        trace_id=trace_id,
        user_id=request.headers.get("X-Gateway-User-ID"),
        item_title=item.title
    )

    new_item = await create_item_in_db(item)

    logger.info(
        "item.created",
        trace_id=trace_id,
        item_id=new_item.id
    )

    return new_item
```

## Best Practices

### 1. Design for Gateway Integration

#### ✓ Do:

- Use standard HTTP status codes
- Return consistent error formats
- Implement health check endpoint
- Support idempotency where applicable
- Use structured logging

#### ✗ Don't:

- Implement your own authentication (use gateway's)
- Perform rate limiting (gateway handles it)
- Log raw request/response bodies (sensitive data)

## 2. Path Design

```
# Good: Clear, RESTful paths
/api/v1/content/items
/api/v1/content/items/{id}
/api/v1/content/categories

# Bad: Inconsistent, unclear
/getItems
/api/items-list
/content_api/v1/get_item_by_id
```

## 3. Versioning

Include version in path:

```
/api/v1/your-service/items # Version 1
/api/v2/your-service/items # Version 2 (breaking changes)
```

## 4. Performance

- Keep responses under 5MB
- Respond to health checks in <100ms
- Set appropriate timeouts
- Use async/await for I/O operations

## 5. Security

- Never return sensitive data in errors
  - Validate all input (even from gateway)
  - Use HTTPS in production
  - Rotate secrets regularly
-

# Examples

## Complete Service Integration

```
# your_service/main.py
from fastapi import FastAPI, Request, HTTPException
import structlog

app = FastAPI(title="Your Service")
logger = structlog.get_logger()

# Health check
@app.get("/api/v1/health")
async def health():
    return {"status": "healthy", "service": "your-service"}

# Protected endpoint
@app.get("/api/v1/your-service/items")
async def get_items(request: Request):
    # Extract user info from gateway headers
    user_id = request.headers.get("X-Gateway-User-ID")
    trace_id = request.headers.get("X-Gateway-Request-ID")

    logger.info("items.list", trace_id=trace_id, user_id=user_id)

    # Your business logic
    items = await fetch_items_from_db()

    return {
        "items": items,
        "trace_id": trace_id
    }

# Error handling
@app.exception_handler(Exception)
async def global_exception_handler(request: Request, exc: Exception):
    trace_id = request.headers.get("X-Gateway-Request-ID")

    logger.error(
        "unhandled.error",
        trace_id=trace_id,
        error=str(exc),
        path=request.url.path
    )

    return {
        "error": {
            "code": "INTERNAL_ERROR",
            "message": "An internal error occurred",
            "trace_id": trace_id
        }
    }
```

## Register with Gateway

```
# scripts/register_with_gateway.py
import requests
import os

GATEWAY_URL = os.getenv("GATEWAY_URL", "http://gateway:8000")
ADMIN_TOKEN = os.getenv("ADMIN_TOKEN")

def register_service():
    response = requests.post(
        f"{GATEWAY_URL}/api/v1/routes",
        headers={
            "Authorization": f"Bearer {ADMIN_TOKEN}",
            "Content-Type": "application/json"
        },
        json={
            "path": "/api/v1/your-service/*",
            "target_service": "your-service",
            "target_url": "http://your-service:8003",
            "methods": ["GET", "POST", "PUT", "DELETE"],
            "auth_required": True,
            "rate_limit": 100,
            "is_active": True
        }
    )

    response.raise_for_status()
    route = response.json()
    print(f"✓ Service registered with route ID: {route['id']}")
    return route

if __name__ == "__main__":
    register_service()
```

## Docker Compose Integration

```
# docker-compose.yml
version: '3.8'

services:
  gateway:
    image: mission-engadi/gateway:latest
    ports:
      - "8000:8000"
    networks:
      - mission-engadi
    environment:
      - DATABASE_URL=postgresql://...
  your-service:
    build: .
    ports:
      - "8003:8003"
    networks:
      - mission-engadi
    environment:
      - GATEWAY_URL=http://gateway:8000
    depends_on:
      - gateway
  networks:
    mission-engadi:
      driver: bridge
```

## Troubleshooting

### Service Not Receiving Requests

1. Check route configuration:

```
bash
curl http://gateway:8000/api/v1/routes
```

2. Verify service is reachable:

```
bash
curl http://your-service:8003/api/v1/health
```

3. Check gateway logs:

```
bash
docker logs gateway-service
```

### Rate Limiting Issues

1. Check rate limit rules:

```
bash
curl http://gateway:8000/api/v1/rate-limits
```

2. Check rate limit status:

```
bash
```

```
curl "http://gateway:8000/api/v1/rate-limits/check?path=/api/v1/your-service/items&client_id=user123"
```

## Authentication Failures

1. Verify route requires auth:

```
json  
{"auth_required": true}
```

2. Check token validity:

```
bash  
jwt decode <your-token>
```

3. Ensure token is in header:

```
bash  
curl -H "Authorization: Bearer <token>" ...
```

---

## Support

- **Documentation:** [docs.engadi.org](https://docs.engadi.org) (<https://docs.engadi.org>)
- **Issues:** [GitHub Issues](https://github.com/mission-engadi/gateway_service/issues) ([https://github.com/mission-engadi/gateway\\_service/issues](https://github.com/mission-engadi/gateway_service/issues))
- **Email:** support@engadi.org